

# Assignment 5

Due at 11:59pm on November 25.

You may work in pairs or individually for this assignment. Make sure you join a group in Canvas if you are working in pairs. Turn in this assignment as an HTML or PDF file to ELMS. Make sure to include the R Markdown or Quarto file that was used to generate it. Include the GitHub link for the repository containing these files.

Working directory at: <https://github.com/TroyLiuUofM/zeng1-liu2-a5.git>

```
library(censusapi)
```

Warning: package 'censusapi' was built under R version 4.3.3

```
library(tidyverse)
```

Warning: package 'ggplot2' was built under R version 4.3.3

Warning: package 'tibble' was built under R version 4.3.3

Warning: package 'purrr' was built under R version 4.3.3

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.4
vforcats   1.0.0     v stringr   1.5.1
v ggplot2   3.5.2     v tibble    3.3.0
v lubridate 1.9.2     v tidyR    1.3.0
v purrr    1.0.4
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()   masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(magrittr)
```

Warning: package 'magrittr' was built under R version 4.3.3

Attaching package: 'magrittr'

The following object is masked from 'package:purrr':

set\_names

The following object is masked from 'package:tidyverse':

extract

```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

## Exploring ACS Data

In this notebook, we use the Census API to gather data from the American Community Survey (ACS). This requires an access key, which can be obtained here:

[https://api.census.gov/data/key\\_signup.html](https://api.census.gov/data/key_signup.html)

```
cs_key = "55d7b56135bd20c3d94c7f4b56d448aa88b66406" #Troy please don't leak my key, lol
acs_il_c <- getCensus(name = "acs/acss5",
                        vintage = 2016,
                        vars = c("NAME",
                                "B01003_001E",
                                "B19013_001E",
                                "B19301_001E"),
                        region = "county:*",
                        regionin = "state:17",
                        key = cs_key) %>%
  rename(pop = B01003_001E,
         hh_income = B19013_001E,
         income = B19301_001E)
head(acs_il_c)
```

	state	county		NAME	pop	hh_income	income
1	17	067	Hancock County, Illinois	18633	50077	25647	
2	17	063	Grundy County, Illinois	50338	67162	30232	
3	17	091	Kankakee County, Illinois	111493	54697	25111	
4	17	043	DuPage County, Illinois	930514	81521	40547	
5	17	003	Alexander County, Illinois	7051	29071	16067	
6	17	129	Menard County, Illinois	12576	60420	31323	

Pull map data for Illinois into a data frame.

```
il_map <- map_data("county", region = "illinois")
head(il_map)
```

	long	lat	group	order	region	subregion
1	-91.49563	40.21018	1	1	illinois	adams
2	-90.91121	40.19299	1	2	illinois	adams
3	-90.91121	40.19299	1	3	illinois	adams
4	-90.91121	40.10704	1	4	illinois	adams
5	-90.91121	39.83775	1	5	illinois	adams
6	-90.91694	39.75754	1	6	illinois	adams

Join the ACS data with the map data. Note that `il_map` has a column `subregion` which includes county names. We need a corresponding variable in the ACS data to join both data sets. This needs some transformations, among which the function `tolower()` might be useful. Call the joined data `acs_map`.

```
# Pull county name from NAME, e.g. "Cook County, Illinois" -> "cook"
acs_il_c <- acs_il_c %>%
  mutate(
    county_name = str_to_lower(str_replace(str_extract(NAME, "[^,]+ County"), " County", ""))
  )

acs_map <- il_map %>%
  left_join(acs_il_c, by = c("subregion" = "county_name"))

# quick check
acs_map %>% select(subregion, pop, income) %>% head()
```

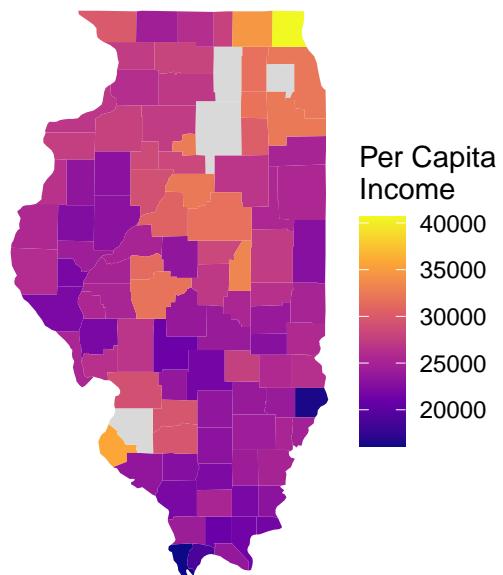
	subregion	pop	income
1	adams	66949	26053
2	adams	66949	26053

```
3     adams 66949  26053
4     adams 66949  26053
5     adams 66949  26053
6     adams 66949  26053
```

After you do this, plot a map of Illinois with Counties colored by per capita income.

```
ggplot(acs_map) +
  geom_polygon(aes(long, lat, group = group, fill = income), color = NA) +
  coord_quickmap() +
  scale_fill_viridis_c(option = "plasma", na.value = "grey85") +
  labs(title = "Illinois Counties: Per Capita Income (ACS 2016)",
       fill = "Per Capita\nIncome") +
  theme_void()
```

Illinois Counties: Per Capita Income (ACS 2016)



## Hierarchical Clustering

We want to find clusters of counties that are similar in their population, average household income and per capita income. First, clean the data so that you have the appropriate variables to use for clustering. Next, create the distance matrix of the cleaned data. This distance matrix can be used to cluster counties, e.g. using the ward method.

Plot the dendrogram to find a reasonable number of clusters. Draw boxes around the clusters of your cluster solution.

```
#Clean data
clust_dat_c <- acs_il_c %>%
  select(subregion = county_name, pop, hh_income, income) %>%
  drop_na()

rownames(clust_dat_c) <- clust_dat_c$subregion
Xc <- scale(clust_dat_c %>% select(pop, hh_income, income)) # standardize

dist_c <- dist(Xc, method = "euclidean")
hc_c   <- hclust(dist_c, method = "ward.D2")

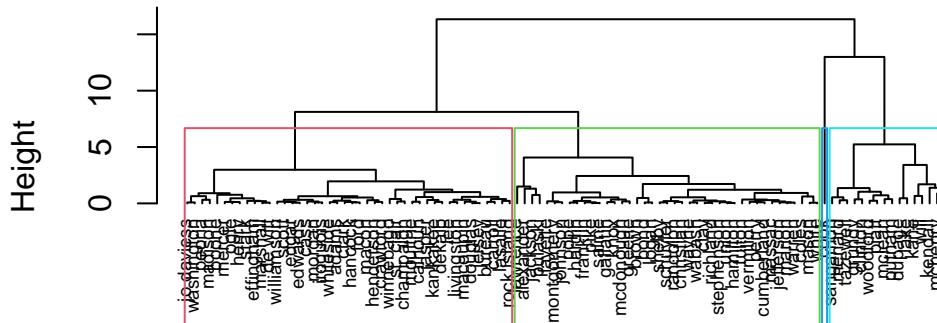
# dendrogram
plot(hc_c, cex = .6, hang = -1, main = "Hierarchical Clustering: IL Counties")

# reasonable number of clusters
Kc <- 4

# cut tree into Kc clusters
clusters_c <- cutree(hc_c, k = Kc)

# draw boxes around clusters in the dendrogram
rect.hclust(hc_c, k = Kc, border = 2:(Kc + 1))
```

## Hierarchical Clustering: IL Counties



```
dist_c  
hclust (*, "ward.D2")
```

```
# check cluster sizes  
table(clusters_c)
```

```
clusters_c  
1 2 3 4  
44 16 41 1
```

Visualize the county clusters on a map. For this task, create a new `acs_map` object that now also includes cluster membership as a new column. This column should be called `cluster`.

## Census Tracts

For the next section we need ACS data on a census tract level. We use the same variables as before.

```
acs_il_t <- getCensus(  
  name      = "acs/acss5",  
  vintage   = 2016,  
  vars      = c("NAME", "B01003_001E", "B19013_001E", "B19301_001E"),  
  region    = "tract:*",  
  regionin= "state:17",
```

```

key      = cs_key
) %>%
  mutate(across(everything(), ~ifelse(. == -666666666, NA, .))) %>% #fun() not working, fixed
  rename(
    pop      = B01003_001E,
    hh_income = B19013_001E,
    income    = B19301_001E
  )

head(acs_il_t)

  state county tract                               NAME   pop
1    17    031 806002 Census Tract 8060.02, Cook County, Illinois 7304
2    17    031 806003 Census Tract 8060.03, Cook County, Illinois 7577
3    17    031 806400    Census Tract 8064, Cook County, Illinois 2684
4    17    031 806501 Census Tract 8065.01, Cook County, Illinois 2590
5    17    031 750600    Census Tract 7506, Cook County, Illinois 3594
6    17    031 310200    Census Tract 3102, Cook County, Illinois 1521
  hh_income income
1      56975  23750
2      53769  25016
3      62750  30154
4      53583  20282
5      40125  18347
6      63250  31403

```

## k-Means

As before, clean our data for clustering census tracts based on population, average household income and per capita income.

```

# Extract county name from NAME
acs_il_t <- acs_il_t %>%
  mutate(
    county = str_extract(NAME, "[^,]+ County"),
    county = str_replace(county, " County", ""),
    county = str_to_lower(county)
  )

clust_dat_t <- acs_il_t %>%
  select(GEO_ID = NAME, county, pop, hh_income, income) %>%

```

```

drop_na()

Xt <- scale(clust_dat_t %>% select(pop, hh_income, income))

```

Since we want to use K Means in this section, we start by determining the optimal number of K that results in Clusters with low within but high between variation. Plot within cluster sums of squares for a range of K (e.g. up to 20).

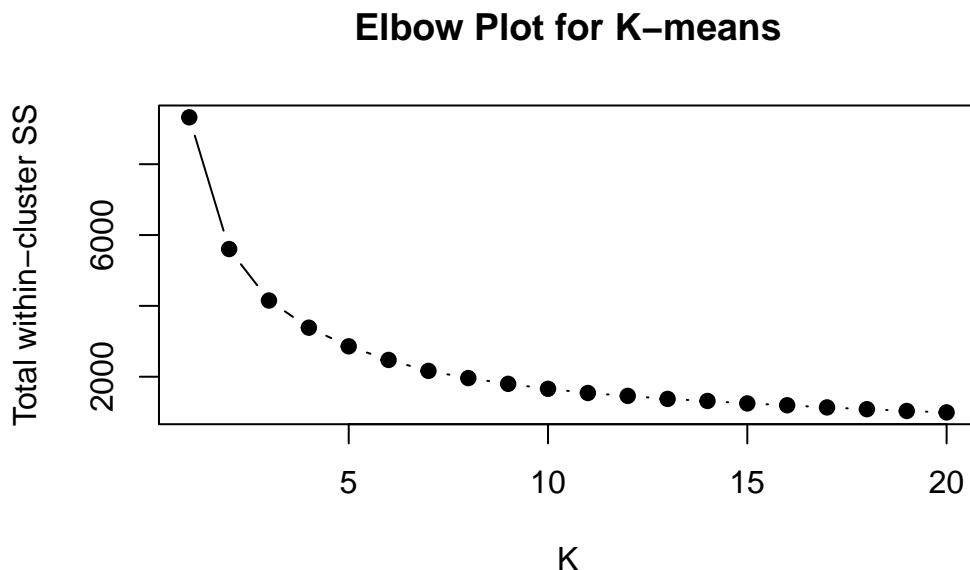
```

#Choose K: elbow (WSS) up to K = 20
kmax <- 20
wss  <- numeric(kmax)

for (k in 1:kmax) {
  km <- kmeans(Xt, centers = k, nstart = 20)
  wss[k] <- km$tot.withinss
}

plot(1:kmax, wss, type = "b", pch = 19, xlab = "K", ylab = "Total within-cluster SS",
      main = "Elbow Plot for K-means")

```



Run `kmeans()` for the optimal number of clusters based on the plot above.

```

K_opt <- 5

set.seed(123) #for reproducability
km_final <- kmeans(Xt, centers = K_opt, nstart = 25)
clust_dat_t$cluster <- km_final$cluster

table(clust_dat_t$cluster)

```

1	2	3	4	5
276	1016	274	744	799

Find the mean population, household income and per capita income grouped by clusters. In addition, display the most frequent county that can be observed within each cluster.

```

clust_summary <- clust_dat_t %>%
  group_by(cluster) %>%
  summarise(
    mean_pop = mean(pop),
    mean_hh_income = mean(hh_income),
    mean_income = mean(income),
    most_freq_county = names(sort(table(county), decreasing = TRUE))[1]
  )

clust_summary

```

	cluster	mean_pop	mean_hh_income	mean_income	most_freq_county
	<int>	<dbl>	<dbl>	<dbl>	<chr>
1	1	3896.	122368.	67665.	"cook"
2	2	2686.	37123.	19778.	"cook"
3	3	7838.	86010.	38154.	"cook"
4	4	5381.	49260.	23275.	"cook"
5	5	3610.	73195.	35913.	"cook"

As you might have seen earlier, it's not always clear which number of clusters is the optimal choice. To automate K Means clustering, program a function based on `kmeans()` that takes K as an argument. You can fix the other arguments, e.g. such that a specific dataset is always used when calling the function.

```

run_kmeans <- function(K) {
  set.seed(123) # for reproducibility
  km <- kmeans(Xt, centers = K, nstart = 25)

  clust_dat_t$cluster <- km$cluster

  clust_summary <- clust_dat_t %>%
    group_by(cluster) %>%
    summarise(
      mean_pop = mean(pop),
      mean_hh_income = mean(hh_income),
      mean_income = mean(income),
      most_freq_county = names(sort(table(county), decreasing = TRUE))[1]
    )

  return(list(
    kmeans_result = km,
    cluster_data = clust_dat_t,
    summary = clust_summary
  ))
}

```

We want to utilize this function to iterate over multiple Ks (e.g., K = 2, ..., 10) and – each time – add the resulting cluster membership as a new variable to our (cleaned) original data frame (`acs_il_t`). There are multiple solutions for this task, e.g. think about the `apply` family or `for` loops.

```

acs_data <- clust_dat_t

# Iterate over K = 2 to 10
for (K in 2:10) {
  set.seed(123)
  km <- kmeans(Xt, centers = K, nstart = 25)

  acs_data[[paste0("cluster_K", K)]] <- km$cluster
}

```

Finally, display the first rows of the updated data set (with multiple cluster columns).

```
head(acs_data)
```

			GEO_ID	county	pop	hh_income	income
1	Census Tract 8060.02, Cook County, Illinois		cook	7304	56975	23750	
2	Census Tract 8060.03, Cook County, Illinois		cook	7577	53769	25016	
3	Census Tract 8064, Cook County, Illinois		cook	2684	62750	30154	
4	Census Tract 8065.01, Cook County, Illinois		cook	2590	53583	20282	
5	Census Tract 7506, Cook County, Illinois		cook	3594	40125	18347	
6	Census Tract 3102, Cook County, Illinois		cook	1521	63250	31403	
	cluster	cluster_K2	cluster_K3	cluster_K4	cluster_K5	cluster_K6	cluster_K7
1	4	2	3	4	4	3	6
2	4	2	3	4	4	3	6
3	5	2	2	1	5	5	3
4	2	2	2	2	2	5	2
5	2	2	2	2	2	4	2
6	5	2	2	1	5	5	3
	cluster_K8	cluster_K9	cluster_K10				
1	3	7	8				
2	3	7	8				
3	2	6	7				
4	2	6	7				
5	6	4	10				
6	2	6	7				